

Untitled

March 28, 2025

```
[9]: import numpy as np
import pandas as pd
from rdkit import Chem
from rdkit.Chem import AllChem
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import classification_report, accuracy_score
```

```
[2]: ds= pd.read_csv("KIBA.csv")
ds =ds.rename(columns={"Ki , Kd and IC50 (KIBA Score)": "binding_affinity"})
ds= ds.sample(n=20000, random_state=42)
print(ds.head())
```

	CHEMBLID	ProteinID	\
58637	CHEMBL1981744	P49760	
7623	CHEMBL185238	Q15759	
28019	CHEMBL1997597	O14757	
72021	CHEMBL1682546	P49760	
85578	CHEMBL2005528	P52333	

	compound_iso_smiles	\
58637	C1=CC(=NC(=C1)NC(=O)NC2=C3C=CC(=CC3=NC=C2)C1)C...	
7623	CC1=CC=CC(=N1)C2=C(C=NN2)C3=NC4=C(C=C3)N=CC=C4	
28019	CCS(=O)(=O)NC1=CC=C(C=C1)C2=CC3=C(C=C2)NN=C3N	
72021	C1=CC=C(C=C1)CNC2=NC(=CS2)C3=CC4=C(C=C3)NN=C4	
85578	CCOC(=O)C1=CC2=C(C=C1)NC3=C2CCNC3=O	

	target_sequence	binding_affinity
58637	MPHPRRYHSSERGRGSYREHYRSRKHKRRRSRSWSSSSDRTRRRR...	11.500000
7623	MSGPRAGFYRQELNKTVWEVPQRLQGLRPVGSGAYGSVC SAYDARL...	10.895880
28019	MAVPFVEDWDLVQTLGEGAYGEVQLAVNRVTEEAVAVKIVDMKRAV...	11.200000
72021	MPHPRRYHSSERGRGSYREHYRSRKHKRRRSRSWSSSSDRTRRRR...	11.800000
85578	MAPPSEETPLIPQRSCSLSTEAGALHVLLPARGPGPPQRLSFSFG...	11.900001

```
[3]: ds= ds.dropna()
print(ds.isna().sum()) # This should print 0 for all columns
```

CHEMBLID	0
----------	---

```
ProteinID          0
compound_iso_smiles 0
target_sequence    0
binding_affinity    0
dtype: int64
```

0.1 Preprocessing the Dataset

```
[4]: #1) for kiba score
threshold=np.mean(ds['binding_affinity'])
ds['target'] = (ds['binding_affinity'] > threshold).astype(int)
print(ds.head())
```

	CHEMBLID	ProteinID	\	compound_iso_smiles	\	target_sequence	binding_affinity	\	target
58637	CHEMBL1981744	P49760		C1=CC(=NC(=C1)NC(=O)NC2=C3C=CC(=CC3=NC=C2)C1)C...		MPHPRRYHSSERGSRGSYREHYRSRKHKRRRSRSWSSSSDRTRRRR...	11.500000		0
7623	CHEMBL185238	Q15759		CC1=CC=CC(=N1)C2=C(C=NN2)C3=NC4=C(C=C3)N=CC=C4		MSGPRAGFYRQELNKTVEVPQRLQGLRPVGSGAYGSVCAYDARL...	10.895880		0
28019	CHEMBL1997597	O14757		CCS(=O)(=O)NC1=CC=C(C=C1)C2=CC3=C(C=C2)NN=C3N		MAVPFVEDWDLVQTLGEGAYGEVQLAVNRVTEEAVAVKIVDMKRAV...	11.200000		0
72021	CHEMBL1682546	P49760		C1=CC=C(C=C1)CNC2=NC(=CS2)C3=CC4=C(C=C3)NN=C4		MPHPRRYHSSERGSRGSYREHYRSRKHKRRRSRSWSSSSDRTRRRR...	11.800000		1
85578	CHEMBL2005528	P52333		CCOC(=O)C1=CC2=C(C=C1)NC3=C2CCNC3=O		MAPPSEETPLIPQRSCSLSTEAGALHVLLPARGPGPPQRLSFSFG...	11.900001		1

```
[5]: def smiles_to_fingerprint(smiles, radius=2, n_bits=1024):
    molecule = Chem.MolFromSmiles(smiles)
    if molecule is not None:
        # Use AllChem.GetMorganFingerprintAsBitVect to generate a fingerprint
        fingerprint = AllChem.GetMorganFingerprintAsBitVect(molecule,
        ↪radius=radius, nBits=n_bits)
        return list(fingerprint)
```

```

else:
    return [0] * n_bits # Return a zero vector if the SMILES is invalid

# Apply the function to generate the fingerprints for each compound in the
↳ dataset
smiles_fingerprints = ds['compound_iso_smiles'].apply(smiles_to_fingerprint)

# Instead of creating individual columns for each fingerprint bit, create a
↳ single vector column
ds['compound_iso_smiles'] = smiles_fingerprints

# Now the 'fingerprints' column contains the fingerprint vectors as a single
↳ list in each row
print(ds.head())

```

	CHEMBLID	ProteinID	\
58637	CHEMBL1981744	P49760	
7623	CHEMBL185238	Q15759	
28019	CHEMBL1997597	O14757	
72021	CHEMBL1682546	P49760	
85578	CHEMBL2005528	P52333	

	compound_iso_smiles	\
58637	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
7623	[0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
28019	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
72021	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
85578	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	

	target_sequence	binding_affinity	\
58637	MPHPRRYHSSERGRSGSYREHYRSRKHKRRRSRSWSSSSDRTRRRR...	11.500000	
7623	MSGPRAGFYRQELNKTVEVPQRLQGLRPVGSGAYGSVC SAYDARL...	10.895880	
28019	MAVPFVEDWDLVQTLGEGAYGEVQLAVNRVTEEAVAVKIVDMKRAV...	11.200000	
72021	MPHPRRYHSSERGRSGSYREHYRSRKHKRRRSRSWSSSSDRTRRRR...	11.800000	
85578	MAPPSEETPLIPQRSCSLSTEAGALHVLLPARGPGPPQRLSFSFG...	11.900001	

	target
58637	0
7623	0
28019	0
72021	1
85578	1

```

[6]: def strip_sequence(seq, length=50):
    # Truncate to the first 'length' characters
    return seq[:length].ljust(length, 'X') # padding with 'X' if the sequence
↳ is shorter than 50

```

```
# Apply the function to the protein sequence column
ds['target_sequence'] = ds['target_sequence'].apply(strip_sequence)

# Print the updated DataFrame
print(ds.head())
```

	CHEMBLID	ProteinID	\
58637	CHEMBL1981744	P49760	
7623	CHEMBL185238	Q15759	
28019	CHEMBL1997597	O14757	
72021	CHEMBL1682546	P49760	
85578	CHEMBL2005528	P52333	

	compound_iso_smiles	\
58637	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
7623	[0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
28019	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
72021	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
85578	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	

	target_sequence	binding_affinity	\
58637	MPHPRRYHSSERGSRGSYREHYRSRKHKRRRSRSWSSSSDRTRRRR...	11.500000	
7623	MSGPRAGFYRQELNKTWVEVPQRLQGLRPVGSGAYGSVCAYDARL...	10.895880	
28019	MAVPFVEDWDLVQTLGEGAYGEVQLAVNRVTEEAVAVKIVDMKRAV...	11.200000	
72021	MPHPRRYHSSERGSRGSYREHYRSRKHKRRRSRSWSSSSDRTRRRR...	11.800000	
85578	MAPPSEETPLIPQRSCSLSTEAGALHVLLPARGPGPPQRLSFSFG...	11.900001	

	target
58637	0
7623	0
28019	0
72021	1
85578	1

```
[7]: # Define the standard 20 amino acids
amino_acids = ['A', 'R', 'N', 'D', 'C', 'Q', 'E', 'G', 'H', 'I', 'L', 'K', 'M', 'F', 'P', 'S', 'T', 'W', 'Y', 'V']

# Function to perform one-hot encoding for a single protein sequence
def one_hot_encode(sequence, amino_acids):
    # Initialize a list to hold the one-hot encoded vectors
    one_hot_encoded = []

    # Iterate over each character in the sequence
    for aa in sequence:
```

```

    # Create a binary vector of length 20, where 1 represents the presence
    ↪ of that amino acid
    encoding = [1 if aa == amino_acid else 0 for amino_acid in amino_acids]
    one_hot_encoded.append(encoding)

    # Return the list of one-hot encoded vectors as a numpy array
    return np.array(one_hot_encoded)

# Sample KIBA dataset (Replace this with your actual `ds` DataFrame)
# For demonstration purposes, I'll create a small subset.
# Replace the following with your actual KIBA dataset (20,000 rows with
    ↪ 'target_sequence' column)

# Apply one-hot encoding to the 'target_sequence' column
ds['target_sequence'] = ds['target_sequence'].apply(lambda seq:
    ↪ one_hot_encode(seq, amino_acids))

# Check the result
print(ds['target_sequence'][0].shape) # Output the shape of the first one-hot
    ↪ encoded sequence

# Display the DataFrame with one-hot encoded sequences (first few rows)
print(ds.head())
# Save the DataFrame to a CSV file
ds.to_csv('output_dataset.csv', index=False) # 'index=False' prevents pandas
    ↪ from writing row numbers (index) to the file

```

(50, 20)

	CHEMBLID	ProteinID	\
58637	CHEMBL1981744	P49760	
7623	CHEMBL185238	Q15759	
28019	CHEMBL1997597	O14757	
72021	CHEMBL1682546	P49760	
85578	CHEMBL2005528	P52333	

	compound_iso_smiles	\
58637	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
7623	[0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
28019	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
72021	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	
85578	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	

	target_sequence	binding_affinity	\
58637	[[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,...	11.500000	
7623	[[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,...	10.895880	
28019	[[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,...	11.200000	
72021	[[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,...	11.800000	

85578 [[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,... 11.900001

	target
58637	0
7623	0
28019	0
72021	1
85578	1

```
[8]: from sklearn.decomposition import PCA

X_smiles = np.array(ds['compound_iso_smiles'].tolist()) # Assumes
↳ 'compound_iso_smiles' is a list of bits (0s, 1s)
X_sequence = np.array(ds['target_sequence'].tolist()) # Assumes
↳ 'target_sequence' is a list of 50x20 matrices

# Concatenate the features: flatten target_sequence (50, 20) into a 1D vector
↳ and concatenate with X_smiles
X = np.hstack((X_smiles, X_sequence.reshape(X_sequence.shape[0], -1))) #
↳ Flatten and concatenate

# Extract target labels (binary classification: 0 or 1)
y = ds['target'].values # Binary target (0 or 1)

# Step 6: Train-test split (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# Step 7: Feature scaling (important for SVM)
scaler = StandardScaler()

# Scale the data
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 8: Train the SVM model
svm_model = SVC(kernel='linear', random_state=42) # You can use other kernels
↳ like 'rbf' or 'poly' too

# Train the model on the reduced data
svm_model.fit(X_train_scaled, y_train)

# Step 9: Evaluate the model
y_pred = svm_model.predict(X_test_scaled)

# Print evaluation metrics
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Accuracy: 0.81775

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.88	0.86	2468
1	0.79	0.72	0.75	1532
accuracy			0.82	4000
macro avg	0.81	0.80	0.80	4000
weighted avg	0.82	0.82	0.82	4000

[]:

[]: